

TRADEPIT

A multi-model debate experiment in crypto paper-trading

three models argue · python decides · nobody risks a penny

Experiment dossier and forward roadmap · revision 9

Live since 20 August 2026 · <https://anchorapp100.github.io/tradepit>

Panel: Claude Sonnet 5 · Gemini 3.5 Flash · DeepSeek V3 · \$10,000 paper per arm · one cycle every 4 hours

Plus a six-model bench: nine minds from five labs, one scoreboard

This document describes a research toy. Nothing in it is financial advice; the system trades simulated money by design, and its expected contribution is knowledge, not income.

1 Executive summary

TRADEPIT asks a simple question with a rigorous setup: **do three AI models that must argue their way to unanimous agreement make better trading decisions than any of them alone?** Every four hours, Claude, Gemini and DeepSeek each read the same market data, form independent views in isolation, debate them through assigned adversarial roles, and cast final ballots. Opening a position requires all three ballots to name the same coin and direction; closing requires two of three. (If a panelist is *technically* unreachable - an API outage or quota, not an opinion - the gate falls back to unanimity of the responding pair, never fewer than two; such cycles are flagged, and a panelist who actually voted flat still blocks opens.) Deterministic Python - which the models cannot override - then applies hard risk limits and executes on a simulated \$10,000 book.

The panel races thirteen rival books: nine controls (each model trading solo on identical data, a seeded coin-flip, buy-and-hold Bitcoin, and four rule-based bots with no AI at all - a trend-follower on both 4-hour and 1-hour clocks, plus a buy-the-dip grid and its identical 'with reflexes' twin) and four **observer arms** that read the debate without influencing it - a Forced Hand that must always trade the panel's single strongest idea, two ghost books trading what the pit rejected, and a ghost that replays the same ballots under a looser 2-of-3 gate (measuring whether unanimity is too strict). An **alarm bell** can wake the panel off-schedule when the market moves violently between cycles (section 3.9). A six-model **bench** of shadow panelists - drawn from five AI labs - gives scored blind stances every cycle with zero influence, building the records a future promotion league will act on (section 3.11). A memory layer (panel-only) scores every blind directional call against the next 24 hours of price and feeds the running track records back into the debate, alongside a journal of post-mortem lessons written whenever a trade stops out. Everything - every prompt, reply, ballot, veto, stance and fill - is stored in SQLite; a public dashboard regenerates after every cycle; and the whole system runs headless on a Raspberry Pi, independent of any desktop being awake.

This dossier records how the system works today (sections 2-3), takes positions on the two big context questions - the four-year cycle and news (section 4) - and lays out a staged roadmap of deterministic, operational, visual and analytical additions (sections 5-8), each tagged with the data it needs before it should be switched on. The governing principle throughout: **change one mechanism at a time, version every change, and keep the control arms frozen** - otherwise the experiment stops being one.

2 The hypothesis, honestly stated

Large language models trained on overlapping data share biases: they chase momentum, they like the same narratives, and they are sycophantic - the first confident opinion anchors the rest. Consensus between them is therefore weaker evidence than it looks, and no debate protocol conjures predictive edge out of correlated minds. What a unanimity gate *can* plausibly do is suppress impulsive trades - and in crypto, where fees and overtrading destroy most retail systems, trading less is a real, measurable improvement.

The experiment is built so that three outputs matter more than the profit-and-loss line:


OUTPUT	QUESTION IT ANSWERS
Calibration	When the panel says 70% confidence, does it win 70% of the time? This is the cleanest test of whether the models know what they know.
The value of the gate	Do the trades that failed consensus, or died at the risk layer, perform worse than the ones that executed? If not, the debate is theatre. Measured live since 21 Aug: the dashboard's Gate Value tile is the pit's return minus the Forced Hand's, and the ghost books track the rejects directly.
Panel vs. solo delta	Does the debate beat each model alone, the average of them, and dumb baselines - or does it just average away both the best and worst ideas?

Anti-anchoring is designed in: views are formed blind before anyone sees anyone else, debate roles (bull advocate, bear advocate, risk officer) rotate every cycle so no model owns a permanent temperament, and failure to agree defaults to flat - being out of the market is treated as a position, not an absence of one.

3 How it works today

3.1 Anatomy of a cycle

Runs every 4 hours at candle close, scheduled by cron on the Pi; the alarm bell (section 3.9) can add off-schedule cycles between them. Each cycle is roughly a dozen model calls.



3.2 Market inputs

4-hour candles for BTC, ETH, SOL, XRP and DOGE (Binance, Kraken fallback) with computed context: 24h/72h/7d change, RSI-14, distance from the 20- and 50-period moving averages, annualised realised volatility, weekly range. Plus perpetual funding rates, the Fear & Greed index, and current news headlines (titles only, with an instruction to judge relevance themselves).

3.3 The deterministic risk layer

Code the models cannot argue with. Every proposal passes through it; every veto is logged and displayed.

RULE	SETTING	WHY
Concurrent positions	max 3, one per coin	diversification floor, blast-radius cap
Position size	5-25% of equity	meaningful but survivable
New exposure per cycle	max 25%	no all-in moments
Stop-loss	mandatory, 2-15%	every trade pre-commits to being wrong
Take-profit	optional, 3-60%	let winners run if the panel wants
Daily loss brake	4% in 24h blocks new opens for 6 cycles	cooling-off period
Circuit breaker	15% drawdown flattens the book, halts 7 days	survival trumps opinion
Costs	0.10% fee per side + 0.05% slippage	a round trip costs ~0.3%; overtrading is punished

3.4 The fourteen arms

ARM	ROLE IN THE EXPERIMENT
The Pit	The treatment: 3-way debate + consensus gate + memory layer.

ARM	ROLE IN THE EXPERIMENT
Claude / Gemini / DeepSeek solo	Each model alone on identical data, no debate, no memory - the controls that isolate what the mechanism adds.
Coin flip	Seeded random trades through the same risk layer - the luck floor any real skill must clear.
Buy & hold BTC	The benchmark every crypto strategy must beat to justify existing.
Quant bot, 4h clock	A classic systematic trend-follower with zero AI: EMA20/50 direction, ATR-scaled stops, ~1% equity risk per trade - what most commercial 'AI trading bots' actually are. The conventional benchmark.
Quant bot, 1h clock	The same strategy deciding hourly on its own cron. The 4h/1h pair is a controlled experiment: does a faster clock capture more of the move, or bleed away in fees and noise?
Grid bot	The 'always trading' retail archetype: buys any coin 3% below its 5-day average, takes profit at +3%, repeat. Thrives in chop; its famous failure mode is a trend.
Grid bot, with reflexes	An identical twin of the grid bot that is also woken by the alarm bell (section 3.9), so it can catch a violent dip the moment it happens instead of at the next 4-hour check. Rules otherwise byte-identical - a single-variable test of whether reflexes help a mechanical strategy or just catch falling knives faster.
Forced Hand (observer)	Must trade the panel's single strongest idea every cycle, consensus or not: the highest-confidence ballot proposal, else the highest-conviction blind stance (8% size, default stop). Deterministic - no extra model calls, no randomness. Pit minus Forced Hand = the live price of requiring agreement.
Ghost: rejected ideas (observer)	Executes every proposal that failed consensus, through the same risk layer - did the gate block good trades?
Ghost: risk vetoes (observer)	Executes what the risk layer killed - what did the vetoes cost or save?
Ghost: 2-of-3 gate (observer)	Replays the same ballots but opens on a majority instead of unanimity - added after a diagnostic showed the pit was trading rarely. Free (no extra debate). Answers whether the 3/3 gate is too strict; early read says no - the constraint is conviction, not the vote count (section 9).

The observer arms read the debate's output but never touch it, and **the panelists are never told they exist** - a model that knew its dissenting idea would be traded anyway would ballot differently, contaminating the treatment. Observers went live at cycle #3 (21 Aug); every observer trade records which panelist's conviction seeded it, so credit and blame stay attributable.

3.5 The memory layer (panel-only)

Track records. Every blind long/short stance is stored with the price at the time and scored against price at **four horizons - 24h, 7d, 14d and 30d** - as each one matures, because a human weighs more than one timescale and a mind that reads the day badly may read the fortnight well. The league table shows all four. Once five calls are scored, each debate opens with the running scoreboard - *"claude 48% (12/25) · gemini 62% (8/13) - weigh each other's arguments accordingly; nobody is exempt, including you."* **Lessons journal.** When a panel trade hits its stop, a rotating panelist writes a post-mortem against the original thesis: a concrete lesson in 35 words or fewer, tagged (chased-momentum, stop-too-tight, against-trend, ...). Future debates read the journal. The solo arms receive neither feature, so the panel-vs-solo gap measures debate and memory together. The records are public as well as internal: each debate card on the site shows the speaker's live hit-rate chip next to their name, moved-in-debate badges credit whose blind idea a convert

adopted, and every contested close vote permanently records who backed the exit and who dissented - the raw material for an eventual answer to "whose caution should count for more?" **Reputations are permanent:** one lifelong scored record per mind, accruing wherever it sits - pit or bench - never reset by promotion, relegation or return. Because a career average is nearly immovable at scale, the league will read the permanent record two ways: career (lifetime, the identity) and form (trailing ~30 scored calls, what seat decisions use).

3.6 What is stored

SQLITE TABLE	CONTENTS
messages	every prompt and raw model reply, parsed JSON, latency, parse status - the full debate record, ~12 rows per cycle
cycles	per-cycle market brief, consensus record, disagreements, vetoes, circuit events
trades	every position with entry thesis, stated confidence, stop/TP, fees, exit reason and P&L
equity	mark-to-market for all fourteen arms, every cycle
stances	every blind per-coin call with its later realised outcome and hit flag
lessons	the post-mortem journal

3.7 The analysis lab (shipped)

The dashboard's ANALYSIS tab studies how the agents *behave*, not just how they score. Six panels, each gated behind an honest sample-size threshold that displays "needs N more" rather than drawing brave lines through four dots: **calibration curves** per panelist (stated conviction vs realised 24h hit rate against the perfect-honesty diagonal); **stance ribbons** (each coin's price with a colour strip per panelist showing every blind call - green long, red short, grey flat, brightness = conviction); the **proposal funnel** (proposed blind, kept on own ballot, reached consensus, executed); **pairwise agreement** (how often each pair lands on the same call independently - do two models gang up on the third?); **role capture** (because debate roles rotate deterministically, we get a natural experiment: does the same model propose more longs when cast as bull advocate than as risk officer? if yes, assigned rhetoric is leaking into judgment); and **debate movement** (rolling share of final ballots that differ from blind takes - does the debate change minds, or is it theatre?). Wave-2 panels - the event-ledger board, ghost-book overlays, lesson-tag effectiveness and the event-vs-swing attribution split - light up as their features land (sections 5-6).

Two weeks of running is roughly 84 cycles, 1,000 logged messages and a few megabytes - enough for the first serious analysis pass (section 9). The running cost is effectively zero: the Claude leg rides an existing subscription, Gemini's free tier covers the cadence, and DeepSeek costs pennies per month.

3.8 Where it runs, and version discipline

Since 21 August the whole system lives on a Raspberry Pi 4 on the home network: cron fires a cycle seven minutes after each 4-hour UTC candle close, an @reboot rule catches up after power cuts, the Claude panelist runs through the household subscription's CLI in headless mode, and the Pi deploys the dashboard itself. No desktop needs to be awake, and the failure mode - silence - is covered: every 4-hour cycle pings a healthchecks.io dead-man's switch on success (a distinct signal on crash), so a missed cycle becomes an email within the hour. The Pi is also its own news supply: after the 2026 paywall wave took every packaged crypto news API (CoinDesk's free tier died in May; CryptoPanic went to \$50/week), the system aggregates publisher RSS, Google News per-coin queries and a free Finnhub wire itself - keyless at its core, hourly, coin-tagged, serving both the dashboard's news page and (from v3) the event ledger. Alpha Vantage's free sentiment endpoint is held in reserve for ledger enrichment. The Gemini leg has its own survival story: when

Google retired the old CLI sign-in for individual accounts (June 2026) and the free API tier's daily quota proved too small for alarm-heavy days, the leg moved to the Antigravity CLI on a Google AI Pro subscription - the same login-not-API-key pattern as the Claude leg - with the API key retained as a wired fallback and the quorum rule (section 1) as insurance. Every cycle is also stamped with an 8-character version hash of the exact configuration and prompt revision that produced it - any future change appears as a marked boundary in the data, so "the panel improved" can always be traced to "after which change".

3.9 The alarm bell (stage 1)

A human trader does not check the market only on a schedule - an alert arrives, and then judgment decides whether it matters. Since 22 August the system works the same way: a 15-minute watcher compares every coin against the last decision point, and a move beyond $\pm 3\%$ convenes an off-schedule cycle on fresh 1-hour candles. The brief opens by stating exactly why the panel was woken, then adds one line of framing: *an alert is information, not an instruction to trade*. Wakes are rationed - at least two hours apart, at most five per day (raised from three after day one spent its whole allowance by mid-afternoon; the cap is a one-line switch, version-stamped like any other configuration change). The solo arms are woken too, so panel-vs-solo stays a fair comparison of judgment rather than of alarm clocks; the rule-based bots are deliberately never woken - their clock is part of their strategy identity - except the grid's 'with reflexes' twin, which exists precisely to test that variable. Alarm cycles are flagged in the data, so alerted-vs-scheduled decision quality is a measurable comparison. Day one produced four genuine wakes - three downside flushes and one upside spike - and the panel stayed flat through three of them.

3.10 Case study: the first trade

Thirteen consecutive cycles ended flat before the first consensus formed - and it took an alarm to do it. On 22 August a $\pm 3\%$ wake found BTC holding above its 50-period average while the rest of the market flushed; the bull advocate proposed the relative-strength long, the bear advocate conceded on the record, the risk officer joined, and the first 3/3 in the experiment's history executed: \$1,500 long BTC at confidence 61, 4% stop. One cycle later the roles had rotated - and the same model that had proposed the trade as bull advocate proposed closing it as bear advocate. The close carried 2-of-3 (closes are deliberately easier than opens); the dissenter's objection - "*closing after one candle is overtrading*" - was overruled, and the round trip finished at -\$4.98, of which \$3.00 was fees.

Five hours of trade bought three findings. **Role capture got its clearest evidence yet:** entry and exit tracked the proposer's assigned costume, not a change in the facts. **The exit violated the entry's own invalidation:** the thesis said 'uptrend structurally intact' and nothing in the exit argued otherwise - 'position already flat/negative' after one candle is the disposition reflex the debate exists to prevent. And **a chosen loss produced no journal entry**, because post-mortems only fire on stop-outs. All three feed the next revision: the rev-3 prompts force close votes to name which part of the entry thesis broke, post-mortems extend to any losing close, and dissent attribution (section 3.5) shipped the same day.

3.11 The bench and the league table

Since 23 August, six shadow panelists give one blind, scored stance per cycle - same briefs, same 24-hour grading, zero influence, and the sitting panel is never told they exist. Two wings: the **reasoning tier** - heavier variants of the incumbents themselves (Claude Opus 5 on the same Claude plan, Gemini 3.7 Flash at high reasoning through the same Antigravity login, DeepSeek Reasoner on the same API key) - three paired sibling contests asking whether more reasoning means better calls; and the **challenger wing** - Kimi K3, Grok 4.6 and GPT-5.x, each through its own consumer-subscription CLI (kimi-cli, Grok Build, Codex). Nine minds from five AI labs now read the same tape every cycle, and the dashboard's league table ranks them all by the same blind test, debaters highlighted. The visibility asymmetry is three-layered and deliberate: debaters see each other's records (that is the memory treatment), debaters do not know the bench exists (survival pressure is applied later, on purpose, versioned), and the bench does not know it is competing (its reads stay honest).

Succession mechanics are already agreed and published on the roadmap: promotion births a fresh \$10,000 solo control; relegation flattens and freezes the old one as a visible museum piece; a returning member un-freezes its own career card, and its equity line carries an honest break for the exile - segmented, never bridged. The bench's purpose is simple: when the league's gates open at ~180 cycles, selection can act on records, not promises.

4 Two context questions: the cycle and the news

4.1 The four-year cycle - lens, not trigger

Bitcoin's halving rhythm produces one of the most discussed patterns in the asset's history, and the panel is currently blind to it - its brief spans seventy days of candles. The record, with the current cycle marked to August 2026:

CYCLE	HALVING	TOP	HALVING TO TOP	TOP TO BOTTOM	MAX DRAWDOWN
1	Nov 2012	Dec 2013 (~\$1,150)	~12.5 mo	~13.5 mo (Jan 2015)	-87%
2	Jul 2016	Dec 2017 (\$19.7k)	~17.4 mo	~12.3 mo (Dec 2018)	-84%
3	May 2020	Nov 2021 (\$69k)	~18.0 mo	~12.4 mo (Nov 2022)	-77%
4	Apr 2024	Oct 2025 (\$126.2k)	~17.6 mo	in progress	-54% so far (Jul 2026)

Three consecutive tops inside a two-week band around 17.5-18 months post-halving is striking - the October 2025 top landed on schedule even while consensus declared the cycle dead. Bottoms rhyme at 12-13.5 months after the top, which frames a historical bottom window through late 2026. And yet: the supply-shock mechanism weakens every cycle (new issuance is now a rounding error against daily volume), the pattern plausibly rides the global liquidity cycle plus reflexive front-running (the first-ever pre-halving all-time high, March 2024, is what a front-run pattern looks like), and $n = 4$ observations would pass no statistical bar anywhere else.

The design conclusion: the **cycle clock** enters the system as deterministic FACTS - months since halving, months since the top, drawdown from the high, the table above, an explicit "four data points, treat as prior" caveat - and never as a conclusion. Phrasing anchors language models hard: a brief that says "the bottom window is open" would turn three debaters into one script. The facts go in (inside the daily regime memo, O3); the interpretation stays contested - that contest is the product.

4.2 News, precedents, and the event ledger

The panel currently sees raw headlines with no memory of what similar news did to price. The **event ledger** (O5) fixes this with event-study discipline. Four traps shape the design. **Attribution is narrative:** media attaches the nearest headline to any move - so the ledger tags ALL salient headlines and scores them mechanically, never only the ones that "worked". **News is priced in faster than a 4-hour cycle:** moves are measured from the first decision point after the system saw the headline - the only window it could ever trade - so the ledger answers the actionable question: does the move continue, fade, or reverse from where we could act? **Scheduled events get front-run:** announcements pop; the event day itself often fades ("buy the rumour, sell the news") - so the ledger distinguishes announcement from occurrence. **Samples start tiny:** every precedent line carries its n , and a risk officer sneering at $n=3$ is the system working as intended.

Learned precedents then pair naturally with **time-boxed exits** (D6): the panel proposes not just a trade but its horizon - "long on this catalyst, exit in six cycles, because measured drift historically ends there" - and the broker enforces the schedule exactly as it enforces stops, with no model awake at exit time. Pre-commitment deletes the disposition effect, every event trade becomes a clean catalyst-horizon-outcome unit, and each one feeds the ledger that taught it - a closed measurement loop.

4.3 The prediction testbed (NEWSTEST)

Before any news analysis is allowed near the panel, the method itself must earn a track record - so a sister experiment runs it walk-forward. A news article goes in; two blinded analysts (Gemini through a tool-less API

call, Grok through a CLI with search disabled - blinding that was *engineered and tested*, not requested: the Antigravity CLI was found fetching live prices even in sandbox mode and is banned from this work) each lock a timestamped prediction: direction at 24h and 7d, expected magnitude, confidence, and what the story is primarily about. Only after locking is the real price revealed and the call graded. Contamination is classed by the calendar alone - cases dated before any analyst's training are quarantined as hindsight-risk and never mixed with clean cases - because the models' self-reported cutoffs proved unreliable in both directions (one claims a 2025 cutoff it does not have; the other confesses overlap for dates after its own release). Since 24 August an autonomous collector completes the loop: it watches the self-aggregated wire, severity-scores fresh stories, picks at most six a day, asks both analysts, grades matured windows, and publishes the ledger to the public site - no human in the loop. Early casework already justifies the layered design: the keyword collector once picked a recycled multi-day story, and the blinded reader caught it ('stale Aug 20 filing'). The graded case library doubles as seed evidence for the event ledger (4.2), and the analysts' clean-case scoreboard will decide whether their opinions ever accompany the statistics: the pit receives receipts, never hearsay.

5 Roadmap - deterministic additions

Pure code: no model judgement involved, so these change the game board rather than the players. They are the safest additions because their effects can be attributed exactly.

D1 Counterfactual ghost books

SHIPPED - DEEPENS WITH DATA

What it is. Two shadow portfolios that paper-trade the trades the system rejected: one executes every proposal that failed consensus, the other every trade the risk layer vetoed. **What it does.** Marks the ghosts to market alongside the real arms and plots their equity next to the Pit's. **Expected contribution.** Measures the mechanism itself. If the no-consensus ghost outperforms the Pit, the gate is destroying value; if it bleeds, the gate is earning its keep. The single most informative addition on this roadmap, and it costs nothing but bookkeeping.

D2 Configuration versioning

SHIPPED - DEEPENS WITH DATA

What it is. A version stamp (config hash + prompt revision) written into every cycle row. **What it does.** Makes every future change - a new prompt, a new feed, a new rule - attributable to a date and version in the data. **Expected contribution.** Pure hygiene, but it is the difference between "the panel improved" and "the panel improved after v7". Prerequisite for every other experiment on this list.

D3 Volatility-scaled stops and sizing

NEEDS ~2 WEEKS OF DATA

What it is. Stops set as a multiple of each coin's recent 4h volatility (ATR-style) instead of a flat percentage; position size shrinks as volatility grows so every trade risks roughly equal equity. **What it does.** Replaces the fixed 2-15% stop band with regime-aware distances; DOGE and BTC stop being treated as equally wild. **Expected contribution.** Noise stop-outs are the classic slow bleed. Two weeks of stop-out statistics will show whether stops are dying inside ordinary candle noise - if so, this is the fix, and the lessons journal will already have been complaining about it.

D4 Confidence-weighted sizing

NEEDS ~2 WEEKS OF DATA

What it is. Position size becomes a function of the panel's mean stated confidence rather than the minimum proposed size. **What it does.** A 90-confidence unanimous trade deploys more than a grudging 60. **Expected contribution.** Makes calibration economically real: if stated confidence carries information, returns improve; if it does not, this amplifies the miscalibration until it is impossible to ignore. Either result is a finding. Gate on the first calibration table having data.

D5 Trailing stops / breakeven move

NEEDS ~2 WEEKS OF DATA

What it is. After a position moves a threshold in its favour, the stop ratchets to entry, then trails. **What it does.** Converts round-trip winners into scratches instead of losses. **Expected contribution.** Cheap asymmetry. Should be added after a baseline period so its effect on win rate and average P&L is visible against the unmodified era.

D6 Time-boxed positions
(panel-scheduled exits)
NEEDS ~2 WEEKS OF DATA

What it is. The action schema gains exit_after_cycles: the panel proposes a trade AND its horizon - typically citing event-ledger precedents ("drift historically ends at ~24h") - and the broker enforces the schedule exactly like stops, no model awake at exit time. A config default TTL backstops trades with no stated horizon. **What it does.** Pre-commits every exit at entry; converts each trade into a clean catalyst-horizon-outcome unit; event trades get labelled as their own trade type for separate attribution. **Expected contribution.** Deletes the disposition effect (models hate closing what they endorsed), makes news-driven trades testable, and - now that stances are scored at four horizons (3.5) - lets the panel commit to the timescale each mind has DEMONSTRATED skill on rather than a guessed one. The mature form: trade the horizon you are proven good at, and quote a distribution of likely outcomes, never a false-precise point return. Stops still apply underneath; the risk layer clamps horizons to 1-18 cycles.

D7 Correlation-aware exposure cap
NEEDS ~2 WEEKS OF DATA

What it is. A cap on net directional exposure across correlated assets (in practice: everything here is BTC-beta), not just per-coin position counts. **What it does.** Three longs in BTC, ETH and SOL count as ~one large bet, and the risk layer treats them as such. **Expected contribution.** The current 3-position cap can still be 75% of equity pointing one direction. This makes diversification real rather than nominal.

D8 Realism upgrades
NEEDS ~1 MONTH / 100+ TRADES

What it is. Funding costs on shorts, spread-aware fills, and per-coin minimum notionals. **What it does.** Narrows the gap between paper results and what an exchange would actually pay. **Expected contribution.** Matters once the numbers start being quoted. Until then the 0.3% round-trip cost already enforces the right instincts.

D9 Kelly-lite sizing from realised edge
NEEDS ~1 MONTH / 100+ TRADES

What it is. Size scales with the arm's measured win rate and payoff ratio, capped well below full Kelly. **What it does.** Sizing driven by evidence of edge rather than enthusiasm. **Expected contribution.** Only meaningful once 100+ closed trades exist; before that it is noise-fitting. The correct long-run endpoint of D4.

6 Roadmap - operational additions

Changes to the AI process itself: who speaks, what they see, how agreement is formed. Higher potential, higher attribution risk - one at a time, versioned via D2.

O1

Reputation-weighted consensus

NEEDS ~1 WEEK OF DATA

What it is. The consensus rule starts consulting the scored track records: e.g. an open passes 2-of-3 when both backers hold a hit rate above 55% over 30+ scored calls, while a panelist below 40% loses veto power over closes. **What it does.** Converts track records from rhetoric (a line in the prompt) into mechanism (votes that count differently). **Expected contribution.** The natural completion of the memory layer. Requires roughly a week of scored stances; the deterministic rule keeps it auditable - no model decides whose vote counts.

O2 Red-team round on formed consensus

NEEDS ~1 WEEK OF DATA

What it is. When a 3/3 open forms, one extra call: the panelist with the lowest stated confidence is ordered to attack the trade and name a fatal flaw if one exists. **What it does.** A confident, specific fatal flaw halves the size or skips the trade (deterministic rule on the reply); a weak objection lets it pass. **Expected contribution.** Directly targets the known failure mode of LLM panels - agreement cascades. Costs one call and only when the panel is about to act, which is exactly when scrutiny is cheapest relative to its value.

O3 Daily regime memo + cycle clock

NEEDS ~1 WEEK OF DATA

What it is. Once a day, a single slower debate produces a one-paragraph regime view - trending or chopping, key levels, what would change the picture - pinned into all six fast cycles that day. The deterministic cycle clock rides inside it: months since halving and since the top, drawdown from the high, the historical rhyme table, and the mandatory "n=4, prior not law" caveat (section 4.1). **What it does.** Separates strategy (daily) from execution (4-hourly), the way trading desks do, and gives the panel cycle context as facts while keeping the interpretation contested. **Expected contribution.** 4-hour cycles currently re-derive the big picture from scratch every time - and debated their first melt-up without knowing they were sitting inside a historical bottom window. A pinned memo fixes both for pennies.

O4 Positioning data feeds

NEEDS ~1 WEEK OF DATA

What it is. Futures open interest, long/short account ratio and taker buy/sell flow (Binance), implied volatility (Deribit DVOL), BTC dominance (CoinGecko). **What it does.** Adds the dimension the panel currently cannot see: what other traders are doing and paying. **Expected contribution.** Funding-plus-OI is the classic crowded-trade tell; a bear advocate with positioning data has real ammunition. Add one feed at a time - a fatter brief also means more anchors and more ways to rationalise anything, so each feed must earn its place through the ghost books and calibration.

O5 Event ledger + precedent memory

NEEDS ~1 WEEK OF DATA

What it is. A cheap model tags each cycle's salient headlines (symbol, category, salience); deterministic code later scores each event's 4h/24h/72h aftermath from the system's first decision point - the same score-later machinery as stances. A digest stage replaces raw titles with three dated bullets. **What it does.** When today's news matches a category, the brief carries receipts: "XRP/political precedents, n=3: +16%, +2%, -4% over the following 24h." Facts with sample sizes, never conclusions - the debate does the interpreting (full design: section 4.2). **Expected contribution.** Turns "this news feels bullish" into measured, timestamped precedent - including the contrarian findings (pops that fade, scheduled events that sell off). Pairs with D6 so learned horizons become scheduled exits.

O6 The league: probation, bench, promotion

NEEDS ~1 MONTH / 100+ TRADES

What it is. A career league in three parts. Probation: the worst scorer receives a truthful, targeted prompt line - hit rate, deadline, consequence - only when replacement is actually on the table. Bench: Kimi, Grok and GPT (and any relegated member) sit as shadow panelists, each logging one scored blind stance per cycle from the same briefs, with zero influence on the debate. Promotion: when the switch date comes, the bench record decides who takes the seat; the relegated model drops to the bench and can earn its way back. **What it does.** The panel's composition itself becomes subject to evidence, and survival pressure becomes a measurable treatment: does a model on notice sharpen or panic? **Expected contribution.** Tests the model-diversity hypothesis directly and keeps the experiment alive as models improve. Strict precondition: an anti-sandbagging rule (a minimum directional-call rate) must land first, or the dominant strategy under pressure is to go permanently flat and freeze your record. The scoring wing went live ahead of schedule on 23 August with six members (section 3.11); this entry now covers only the pressure mechanics - probation notices, promotion windows, relegation - which promote on FORM (trailing window), display CAREER always, and never bluff.

O7 Prompt self-tuning loop

NEEDS ~1 MONTH / 100+ TRADES

What it is. Monthly, a model reads the full analysis report and proposes one specific prompt change with a rationale; a human approves; the version bumps (D2). **What it does.** Closes the meta-learning loop: the experiment starts improving its own harness. **Expected contribution.** Genuinely interesting - and genuinely dangerous without discipline. One change per month, human-gated, fully versioned, reversible. The transcripts make the before/after measurable.

O8 Always-on hosting SHIPPED - DEEPENS WITH DATA

What it is. Done, ahead of schedule: the orchestrator moved to a Raspberry Pi 4 on 21 August (section 3.8), with the desktop schedule retired the same hour so the histories could never fork. **What it does.** Cycles no longer depend on any PC being awake; missed-cycle gaps in the data close. **Expected contribution.** Data quality more than convenience: gaps break 24h stance scoring windows and leave positions unattended past their stops. The healthcheck dead-man's switch is live too: every cycle pings a monitor that emails if the Pi goes quiet.

7 Roadmap - visual additions

The dashboard already shows equity, the consensus strip, the leaderboard, per-coin charts with trade markers, changed-mind badges, full transcripts, the lessons journal - and, as of revision 2, the ANALYSIS lab (section 3.7): calibration curves, stance ribbons, the proposal funnel, pairwise agreement, role capture and debate movement, each behind honest sample-size gates. Still to come:

V1 Calibration curves - full depth

SHIPPED - DEEPENS WITH DATA

What it is. The basic per-panelist honesty curve is live; full depth arrives with data. **What it does.** Shows who is overconfident, who sandbags, and whether 80 ever means 80. **Expected contribution.** The single most character-revealing chart the data can produce. Points render at 3+ scored calls per bucket; the shape becomes trustworthy near ~50 calls per panelist.

V2 Ghost-book overlays

SHIPPED - DEEPENS WITH DATA

What it is. The counterfactual equity curves are live: dashed grey lines beside the real arms, with the Gate Value tile (pit minus Forced Hand) at the top of the overview. **What it does.** "What the vetoed trades would have done" is now a line you can point at. **Expected contribution.** Turns the most important analytical question into the most legible picture on the site. Shipped with D1 in the observer batch.

V3 Exposure timeline

READY NOW

What it is. A band under the equity chart: percent of equity deployed, net long vs short, per cycle. **What it does.** Shows when the pit was brave, when it hid in cash, and how exposure correlated with what happened next. **Expected contribution.** Cheap to build from existing equity and trade data; makes the flat stretches readable as decisions rather than absences.

V4 Debate-dynamics view - full depth

SHIPPED - DEEPENS WITH DATA

What it is. The rolling debate-movement chart and per-cycle changed-mind badges are live; the remaining piece is the per-cycle stance grid with concession arrows and who-moves-first statistics. **What it does.** Makes the social dynamics of the panel - anchoring, stubbornness, capitulation - visible over time. **Expected contribution.** The transcripts already contain this; the view surfaces whether consensus is reasoned convergence or polite surrender.

V5 Per-coin and per-tag analytics

NEEDS ~2 WEEKS OF DATA

What it is. P&L broken down by coin, and lesson-journal tags plotted over time. **What it does.** Reveals whether the pit has a profitable niche (or a nemesis coin), and whether tagged mistakes actually stop recurring after entering the journal. **Expected contribution.** The lessons-tag trend is the direct test of whether the journal changes behaviour - the memory layer's own report card.

V6 Weekly digest card

NEEDS ~1 WEEK OF DATA

What it is. An auto-generated Monday summary at the top of the overview: cycles run, trades, best and worst call, biggest disagreement, quote of the week from the transcripts. **What it does.** A narrative layer over the numbers, written by one cheap model call from the DB. **Expected contribution.** Makes the site worth visiting weekly even when the pit has been sitting on its hands - which, if the design works, is often.

V7 Phone-grade polish

READY NOW

What it is. PWA manifest (installs like an app), social-preview meta tags, small-screen layout passes. **What it does.** The dashboard becomes a proper pocket scoreboard and shares cleanly in chats. **Expected contribution.** Low effort, permanent quality-of-life gain for the primary way the site will actually be checked.

8 Roadmap - analytical additions

**A1 Report v2:
uncertainty and
attribution**
NEEDS ~2 WEEKS OF DATA

What it is. The --report command grows bootstrap confidence intervals on every return figure, regime-split performance (up-weeks vs down-weeks), and a debate-delta line: panel return minus best solo, mean solo, and coin-flip. **What it does.** Every claim arrives with its error bars. **Expected contribution.** The defence against the strongest failure mode of experiments like this: narrating noise. If the intervals overlap, the report says so in plain text.

A2 Transcript mining
NEEDS ~1 MONTH / 100+
TRADES

What it is. Systematic passes over the message log: argument patterns that precede wins vs losses, concession dynamics, first-mover anchoring, phrase-level tells. **What it does.** Turns a thousand debates into behavioural findings about how these models argue under uncertainty. **Expected contribution.** This is where the experiment's most publishable insights likely live - the P&L is small-N, but the linguistic dataset is rich from week one.

A3 Data export
NEEDS ~1 WEEK OF DATA

What it is. One command dumping tidy CSVs (cycles, trades, stances, ballots) for notebook analysis. **What it does.** Frees the analysis from the site's rendering and the report's format. **Expected contribution.** Also the on-ramp for handing the dataset to an AI for the periodic deep-dive reviews.

9 Milestone plan

Work splits into two families with different rules. **Observer additions** (ghost books, Forced Hand, versioning, the ANALYSIS lab) read the experiment without touching the treatment - they shipped immediately, and did (day one, ahead of schedule). **Mechanism additions** change what the panel sees or how consensus works - they wait for a clean baseline and land as marked version boundaries. The alarm bell was pulled forward from this wave at the operator's request and shipped 22 August (section 3.9). The next mechanism wave - "v3" - is **O3 regime memo with the cycle clock**, **O5 event ledger**, **D6 time-boxed exits**, the **rev-3 debate prompts** (synthesis nudge, forced engagement with invalidation levels, and - after the first trade - close votes that must name which part of the entry thesis broke), and **post-mortems extended to losing decision-closes**.

A worked example of the discipline (24 Aug). Three days in, the pit had opened a single trade, and the obvious reform was to loosen the consensus gate from 3-of-3 to 2-of-3. Before changing anything, a diagnostic counted the actual ballots across 31 cycles: in 19 of them (61%) *no* panelist wanted to trade; only twice did as many as two agree on the same open. Dropping to 2-of-3 would have added two trades in three days - and one of those still fails a conviction floor. The threshold was never the constraint; conviction was, and three independent models declining to chase an overbought melt-up is discipline, not paralysis (the Forced Hand, compelled to act, sits last). So the gate was left untouched and a free 2-of-3 ghost added to prove it on the record - while the genuine fix, multi-horizon scoring (3.5), addresses the real suspicion: that grading only at 24h was suppressing conviction. The reflex was to change the mechanism; the evidence said change the measurement. That ordering is the method.

WHEN	DATA ON HAND	WHAT HAPPENS
Day 3	~18 cycles, first scored stances	Plumbing review. Hit rates appear on the leaderboard, the lab's gated panels start unlocking, track-record lines enter the prompts, and the observer books accumulate their first verdicts on the gate.
Week 1	~42 cycles, 30+ scored calls per panelist	The v3 mechanism batch lands (O3+clock, O5, D6). Decision point on O1 (reputation-weighted consensus) and O2 (red-team round).
Week 2	~84 cycles, first closed-trade sample	First full analysis pass over the DB: calibration, gate value via ghost books, panel-vs-solo delta, stop-out autopsy, first event-ledger readouts. The evidence picks the next additions - D3/D4 are the likely candidates.
Month 1	~180 cycles, meaningful trade count	Calibration verdict. Consider O6 (relegation) and the first O7 prompt revision. Publish findings-so-far on the site's About page.
Ongoing	-	One mechanism change at a time. Solo arms stay frozen forever - they are the control group, and the moment they get upgrades the experiment loses its baseline.

10 Limitations and honesty notes

Paper is not production. Shorts carry no funding cost, fills are simplified, and liquidity is assumed. Results are indicative, not exchange-accurate (D8 narrows this). **Small N.** At 4-hour cadence, weeks of running produce dozens - not thousands - of trades; every performance number should be read with A1's error bars once they exist, and until then treated as anecdote. **Correlated panelists.** Three models that agree may simply share a blind spot; the ghost books and solo arms exist precisely because agreement alone proves nothing. **Tool-access asymmetry.** The nine minds reach us through different doors - API calls (structurally tool-less), and agentic CLIs whose web access varies: Grok's is explicitly disabled, Codex runs in a no-network sandbox, but the Antigravity-served Gemini legs cannot be flag-blinded (reply latencies suggest no browsing in practice, but 'can but probably doesn't' is not a control condition), and Kimi's print mode

auto-approves tools. Until sealed, this is a known, documented asymmetry in what the minds could in principle consult. **Prompt sensitivity.** Small wording changes move LLM behaviour; that is why every prompt revision is versioned and rate-limited to one per month. **And the obvious one:** nothing here is financial advice, the money is simulated by design, and the expected return on this project is understanding - how AI models reason, argue, concede and miscalibrate under uncertainty - not income.

TRADEPIT · <https://anchorapp100.github.io/tradepit> · dossier revision 9 · 25 August 2026